

Scalable Synchronization on Shared Memory Multiprocessors

- Paper discusses several algorithms for implementing spinlocks and barriers; includes performance discussion on two types of hardware

Intro

- Why need synchronization constructs in the kernel?
- Can a mutex be sufficient?

The lock in GTTreads

```
int actual_lock(int* mutex) {  
  
    int mutex_val = 1;  
  
    while(mutex_val)  
    {  
        asm("movl $0x1, %%eax;"  
            "movl %1, %%ebx;"  
            "xchg %%eax, (%%ebx);"   
            "movl %%eax, %0;"  
            : "=m" (mutex_val)  
            : "m" (mutex)  
            : "%eax", "%ebx");  
    }  
    return 0;  
}
```

It's spinning!

Can do it differently so it blocks...

Basic sync constructs

- Blocking – may be too heavy
- Busy-waiting
 - Benefits when?
- Built on top of some atomic instructions supported by hardware
 - Different on different platforms
 - test-and-set(<mem>)
 - fetch-and-store(R, <mem>)
 - fetch-and-add(<mem>, <value>)
 - compare-and-swap(<mem>, <cmp-val>, <stor-val>)
 - Execution costly – typically atomic instruction costs more cycles than a regular one

Basic spin-lock

- Using atomic test_and_set operation

```
init      lock = clear;  
lock      while(t&s(lock) == busy); /* spin */  
unlock    lock = clear;
```

- Where is lock?
- Do we have to spin on the shared variable (and the atomic instruction)?
- What is the architecture like?
 - Bus-based vs. Interconnection network
 - Local caches? Cache-coherency?
 - Write-update (Broadcast)
 - Write-invalidate

Algorithm Objectives

- reduce latency
 - how quickly can I get the lock in the absence of competition
- reduce waiting time – delaying the application
 - time between lock is freed and another processor acquires it
- reduce contention / bandwidth consumption
 - other processors need to do useful work, process holding lock needs to complete asap...
 - how to design the waiting scheme to reduce interconnection network contention

Any conflicts here??

Reduce contention approach – Backoff delay

```
type lock = (unlocked, locked)
procedure acquire_lock(L:^lock)
    delay : integer :=1
    while test_and_set(L) = locked
        pause(delay)
        delay:=delay*2
procedure release_lock(L:^lock)
    lock^:=unlocked
```

- Different processors back-off different amount of time when they see lock was busy
- Exponential backoff, static backoff...
- Delay approaches work for bus-based and interconnection network based architectures (“dance-hall”)
- Problem -?

Reduce waiting time/delay

- Can you find out immediately that lock is freed without going to memory location?
- Caches! & cache-coherent architectures
- Idea – spin on locally cached value, act when you see a change
- Test_and_test&set

```
lock      while (lock == busy) spin;
          if (t&s(lock) == busy) goto lock;
```
- What's behaviour like in Write-Update vs. Write-Invalidate?
- What if you don't have hardware supported coherency?
Can all local “lock” copies be changed?

Ticket lock

```
type lock = record
  next_ticket : unsigned integer := 0
  now_serving : unsigned integer := 0

procedure acquire_lock (L : ^lock)
  my_ticket : unsigned integer := fetch_and_increment (&L->next_ticket)
  // returns old value; arithmetic overflow is harmless
  loop
    pause (my_ticket - L->now_serving)
    // consume this many units of time
    // on most machines, subtraction works correctly despite over
  if L->now_serving = my_ticket
    return

procedure release_lock (L : ^lock)
  L->now_serving := L->now_serving + 1
```

What if not supported?


Backoff time good/bad?

private

shared

Can we spin on local and private?

Queuing locks



- a unique ticket for each arriving processor
- get ticket in waiting queue
- go into CS when you have lock
- signal next waiter when you are done
- assume a primitive op `r&inc(mem)`

Anderson's array-based queuing lock

```
type lock = record
  slots : array [0..numprocs -1] of (has_lock, must_wait)
    := (has_lock, must_wait, must_wait, ..., must_wait)
    // each element of slots should lie in a different memory module
    // or cache line
  next_slot : integer := 0

// parameter my_place, below, points to a private variable
// in an enclosing scope
```

```
procedure acquire_lock (L : ^lock, my_place : ^integer)
  my_place^ := fetch_and_increment (&L->next_slot)
    // returns old value
  if my_place^ mod numprocs = 0
    atomic_add (&L->next_slot, -numprocs)
    // avoid problems with overflow; return value ignored
  my_place^ := my_place^ mod numprocs
  repeat while L->slots[my_place^] = must_wait
    L->slots[my_place^] := must_wait

procedure release_lock (L : ^lock, my_place : ^integer)
  L->slots[(my_place^ + 1) mod numprocs] := has_lock
```

again!
approach: replace with
Regular test_and_set
Noone is spinning on next_slot

make sure it's
in different cache
line/Mm module

Graunke-Thakkar's array-based lock

```
type lock = record
  slots : array [0..numprocs -1] of Boolean := true
    // each element of slots should lie in a different memory module
    // or cache line
  tail : record
    who_was_last : ^Boolean := 0
    this_means_locked : Boolean := false
    // this_means_locked is a one-bit quantity.
    // who_was_last points to an element of slots.
    // if all elements lie at even addresses, this tail "record"
    // can be made to fit in one word
  processor private vpid : integer // a unique virtual processor index

procedure acquire_lock (L : ^lock)
  (who_is_ahead_of_me : ^Boolean, what_is_locked : Boolean)
    := fetch_and_store (&L->tail, (&slots[vpid], slots[vpid]))
  repeat while who_is_ahead_of_me^ = what_is_locked

procedure release_lock (L : ^lock)
  L->slots[vpid] := not L->slots[vpid]
```

avoids read&inc

Issues with these locks

- Proc: FIFO ordering
 - Delay, resource contention – good
 - (not in ticket without CCwBroadcast)
 - Though may be dangerous if process in the queue gets context switched!
 - Ticket lock also guarantees FIFO.
- Cons:
 - lock structure size proportional to # of procs.
 - Latency issue, especially if no r&inc in Anderson, ticket

MCS queuing lock

```
type qnode = record
    next : ^qnode
    locked : Boolean
type lock = ^qnode

// parameter I, below, points to a qnode record allocated
// (in an enclosing scope) in shared memory locally-accessible
// to the invoking processor

procedure acquire_lock (L : ^lock, I : ^qnode)
    I->next := nil
    predecessor : ^qnode := fetch_and_store (L, I)
    if predecessor != nil      // queue was non-empty
        I->locked := true
        predecessor->next := I
        repeat while I->locked      // spin

procedure release_lock (L : ^lock, I: ^qnode)
    if I->next = nil      // no known successor
        if compare_and_swap (L, I, nil)
            return
        // compare_and_swap returns true iff it swapped
        repeat while I->next = nil      // spin
    I->next->locked := false
```

No known successor?!

- Time difference between `fetch_and_store` in `acquire_lock` & `predecessor->next = I`
- If we think no one's behind us, we must atomically (1) verify this, and (2) set lock to `NIL`
 - `Compare_and_swap` solves this

No compare_and_swap

```
procedure release_lock (L : ^lock, I : ^qnode)
  if I->next = nil          // no known successor
    old_tail : ^qnode := fetch_and_store (L, nil)
    if old_tail = I        // I really had no successor
      return
    // we have accidentally removed some processor(s) from the queue;
    // we need to put them back
    usurper := fetch_and_store (L, old_tail)
    repeat while I->next = nil          // wait for pointer to victim list
    if usurper != nil
      // somebody got into the queue ahead of our victims
      usurper->next := I->next          // link victims after the last usurper
    else
      I->next->locked := false
  else
    I->next->locked := false
```

- Doable – but now 2 atomic operations – they are expensive!

Performance comparisons

- BBN Butterfly
 - Shared mm multiprocessor, up to 256 processor nodes, \log_4 -depth interconnection network
 - Atomic ops `fetch_and_clear_then_add` and `fetch_and_clear_then_xor`; mask used to implement other ops, including `fetch_and_add`;
 - Atomic ops very expensive compared to regular operations
- Sequent Symmetry
 - Shared bus multiprocessor, up to 30 nodes
 - Write-invalidate cache-coherency
 - `Fetch_and_store` supported; other logical/arithm ops can be applied atomically but do not return a value (e.g., `incr` vs. `read_and_incr`)
 - Atomics cheaper
- Neither supports `compare_and_swap`
- Single processor #s:

	<code>test_and_set</code>	<code>ticket</code>	Anderson	MCS
Butterfly	34.9 μs	38.7 μs	65.7 μs	71.3 μs
Symmetry	7.0 μs	NA	10.6 μs	9.2 μs




Figure 5: Performance of selected spin locks on the Butterfly (empty critical section)




Figure 6: Performance of spin locks on the Symmetry (empty critical section)




Figure 7: Performance of spin locks on the Symmetry (small critical section).


Network contention on interconnect-based multiprocessors

Busy-wait Lock	Increase in Network Latency Measured From	
	Lock Node	Idle Node
<code>test_and_set</code>	1420%	96%
<code>test_and_set</code> w/ linear backoff	882%	67%
<code>test_and_set</code> w/ exp. backoff	32%	4%
<code>ticket</code>	992%	97%
<code>ticket</code> w/ prop. backoff	53%	8%
Anderson	75%	67%
MCS	4%	2%

- BBN with 60 processors

Barriers

- User in parallel apps, to make sure no process advances beyond a point until all processes in the group reach the barrier



Basic sense-reversing centralized barrier

```
shared count : integer := P
shared sense : Boolean := true
processor private local_sense : Boolean := true

procedure central_barrier
  local_sense := not local_sense  // each processor toggles its own sense
  if fetch_and_decrement (&count) = 1
    count := P
    sense := local_sense          // last processor toggles global sense
  else
    repeat until sense = local_sense
```

- What's the purpose of the sense variable?
- Can we implement barriers without sense reversing?
- Spins on shared sense variable
 - when is this ok?
 - Hint: only one processor is updating
- Does inserting delays, like with spinlocks, help?

Software Combining Tree Barrier

```
type node = record
  k : integer           // fan-in of this node
  count : integer       // initialized to k
  locksense : Boolean   // initially false
  parent : ^node        // pointer to parent node; nil if root

shared nodes : array [0..P-1] of node
  // each element of nodes allocated in a different memory module or cache line
processor private sense : Boolean := true
processor private mynode : ^node  // my group's leaf in the combining tree

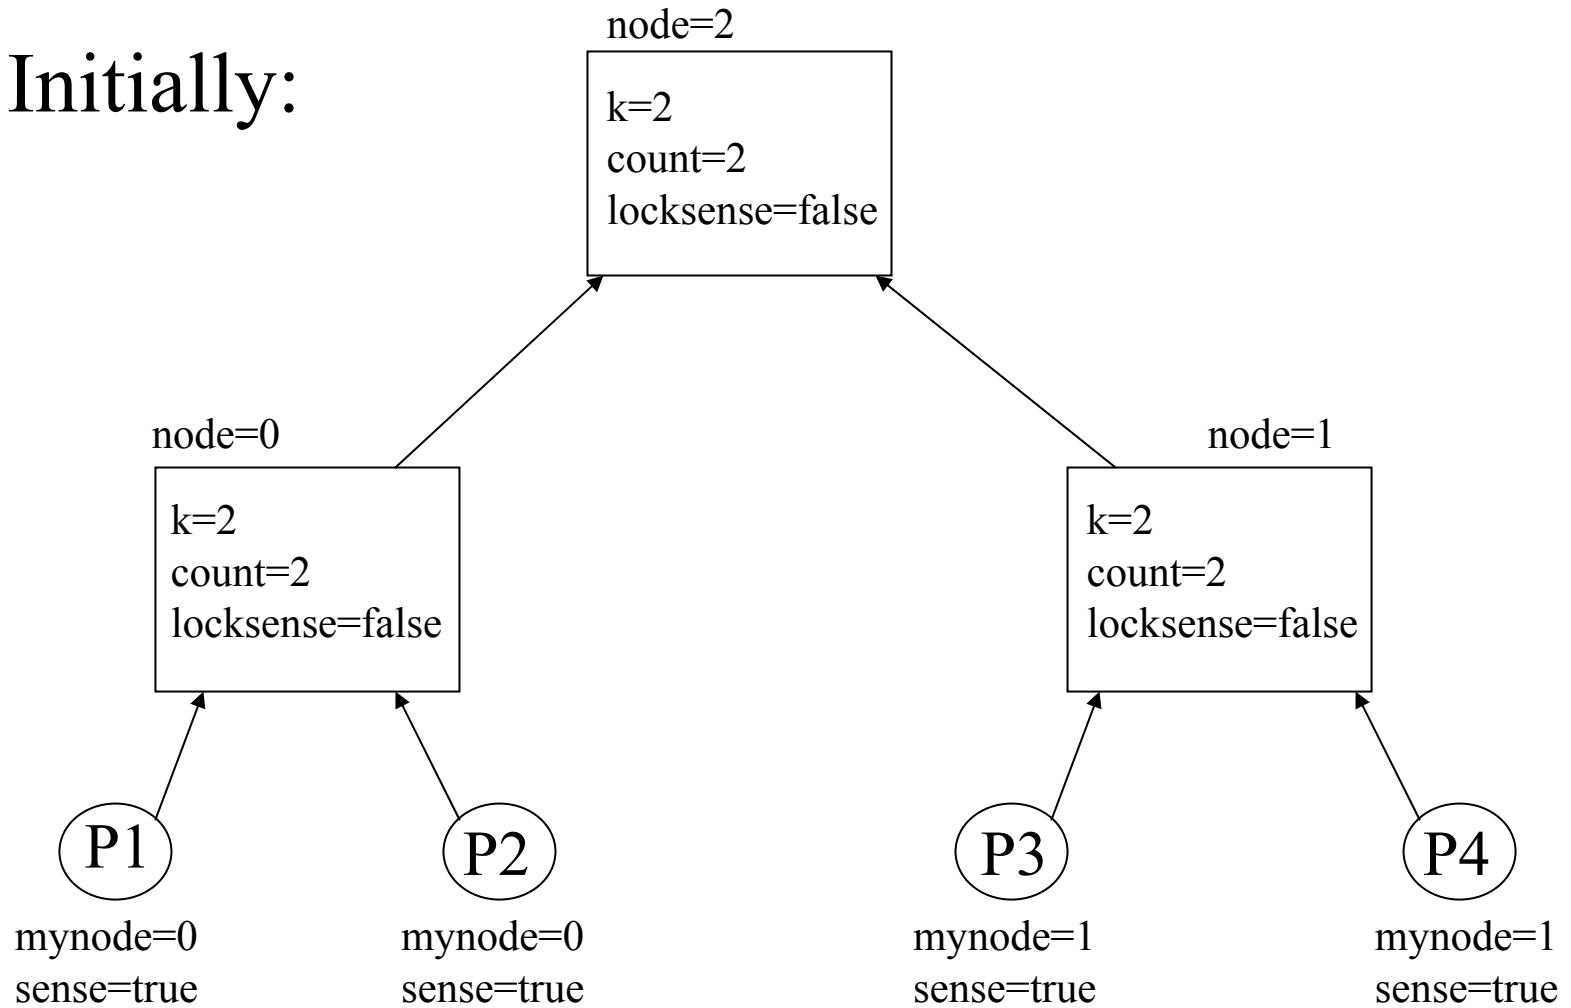
procedure combining_barrier
  combining_barrier_aux (mynode)    // join the barrier
  sense := not sense               // for next barrier

procedure combining_barrier_aux (nodepointer : ^node)
  with nodepointer^ do
    if fetch_and_decrement (&count) = 1    // last one to reach this node
      if parent != nil
        combining_barrier_aux (parent)
      count := k                          // prepare for next barrier
      locksense := not locksense          // release waiting processors
    repeat until locksense = sense
```


(illustration – a tree wouldn't necessarily look like this binary tree)

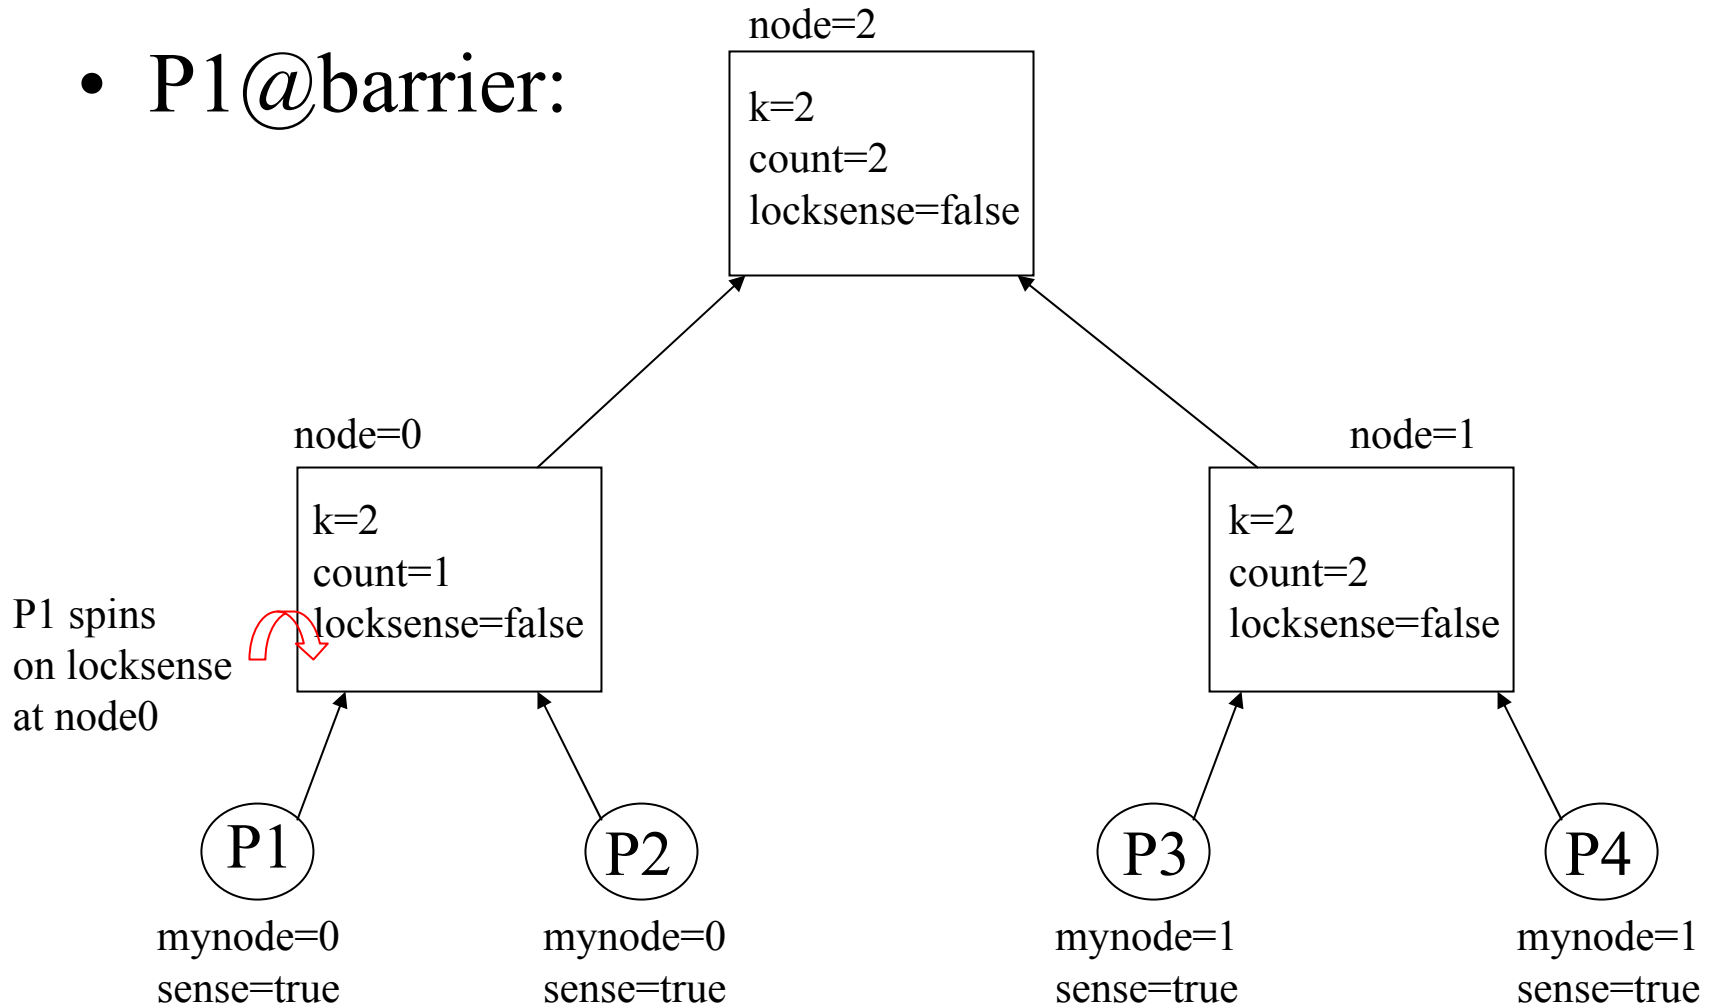
Software Combining Tree Barrier

- Initially:



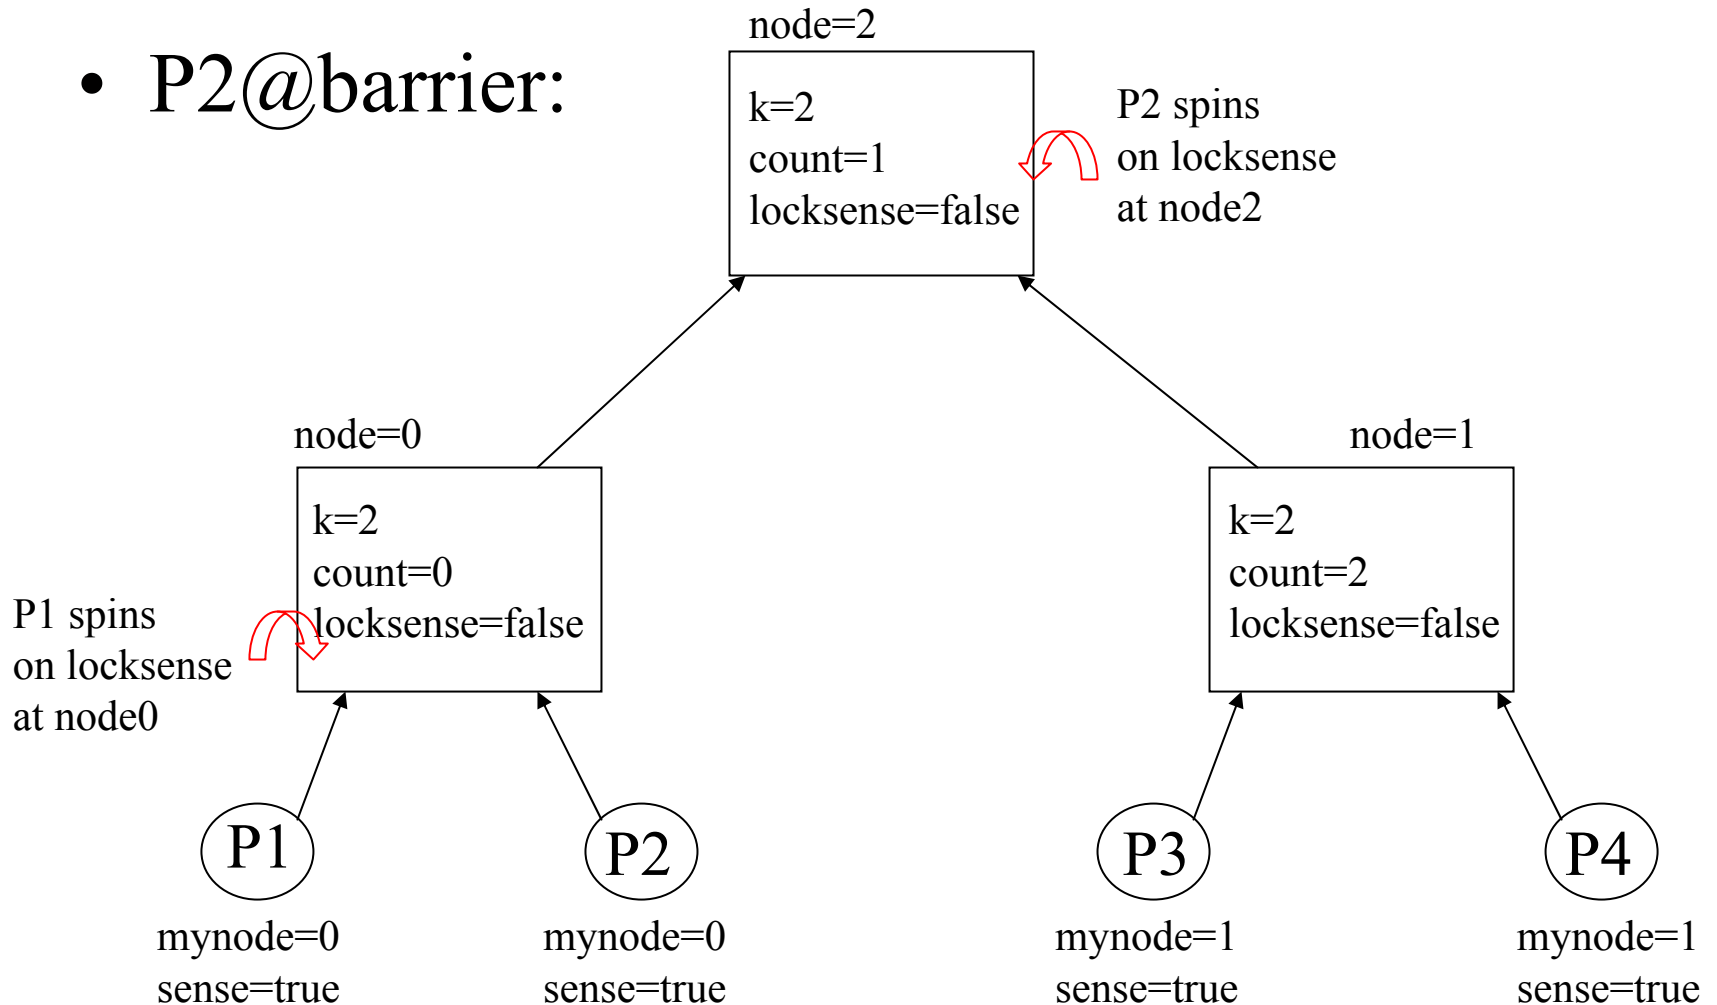
Software Combining Tree Barrier

- P1@barrier:



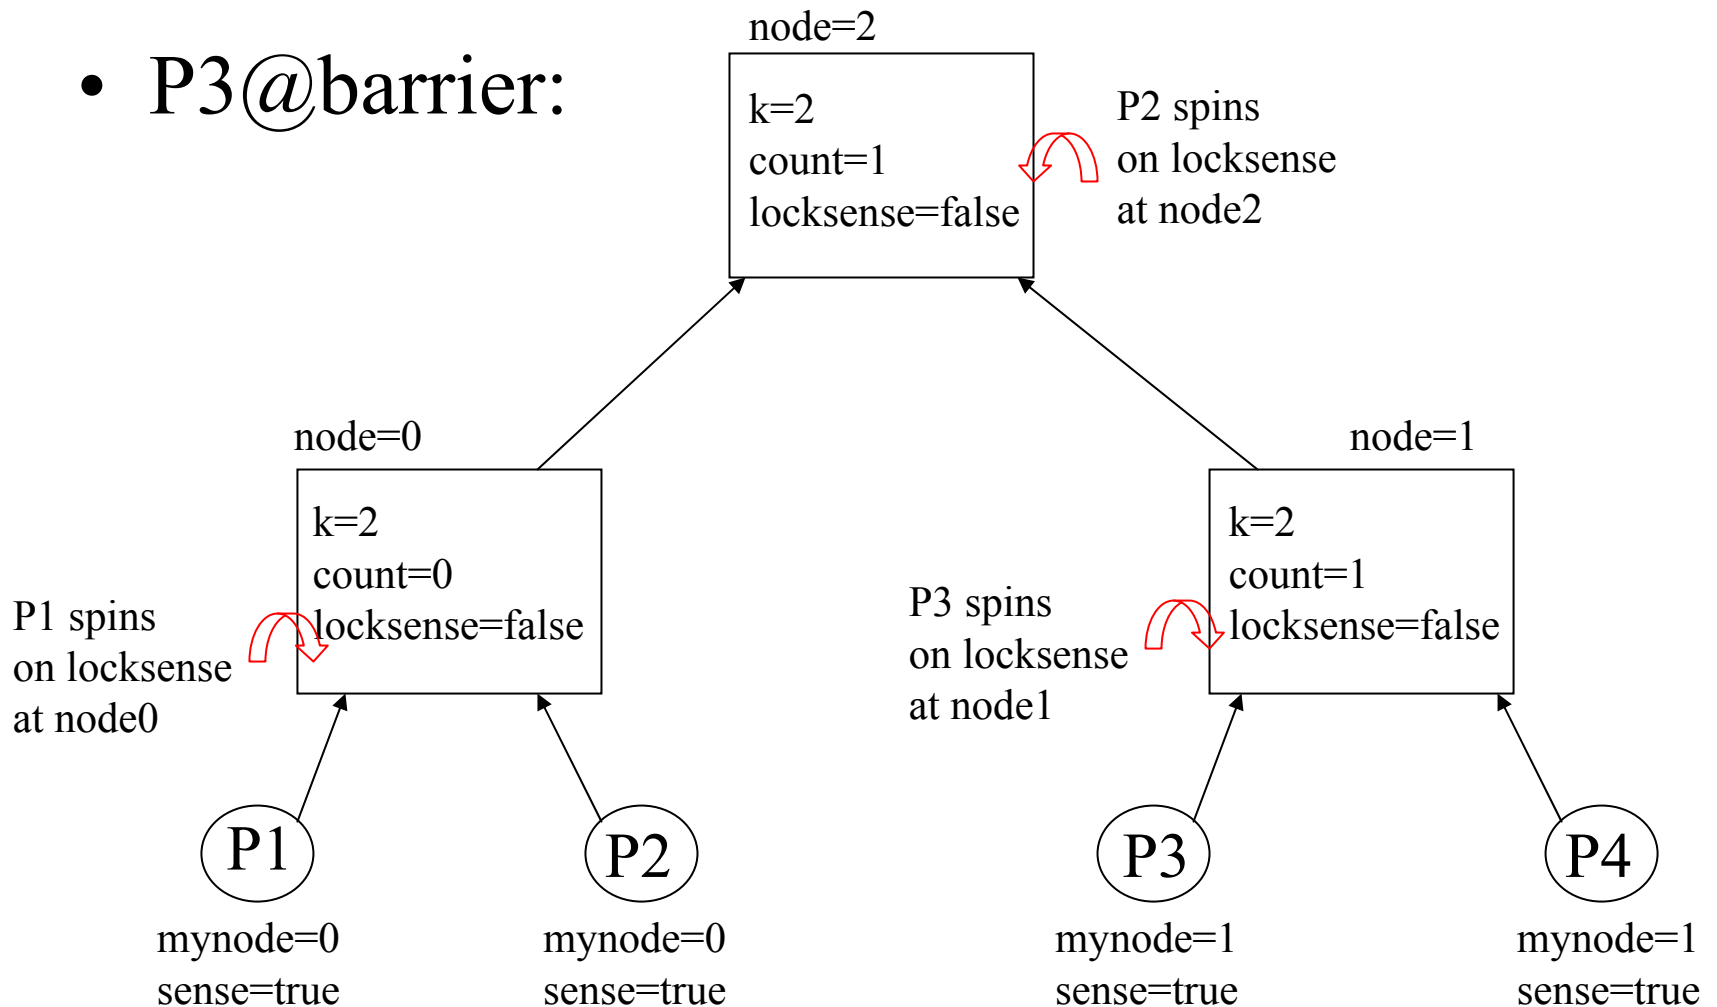
Software Combining Tree Barrier

- P2@barrier:



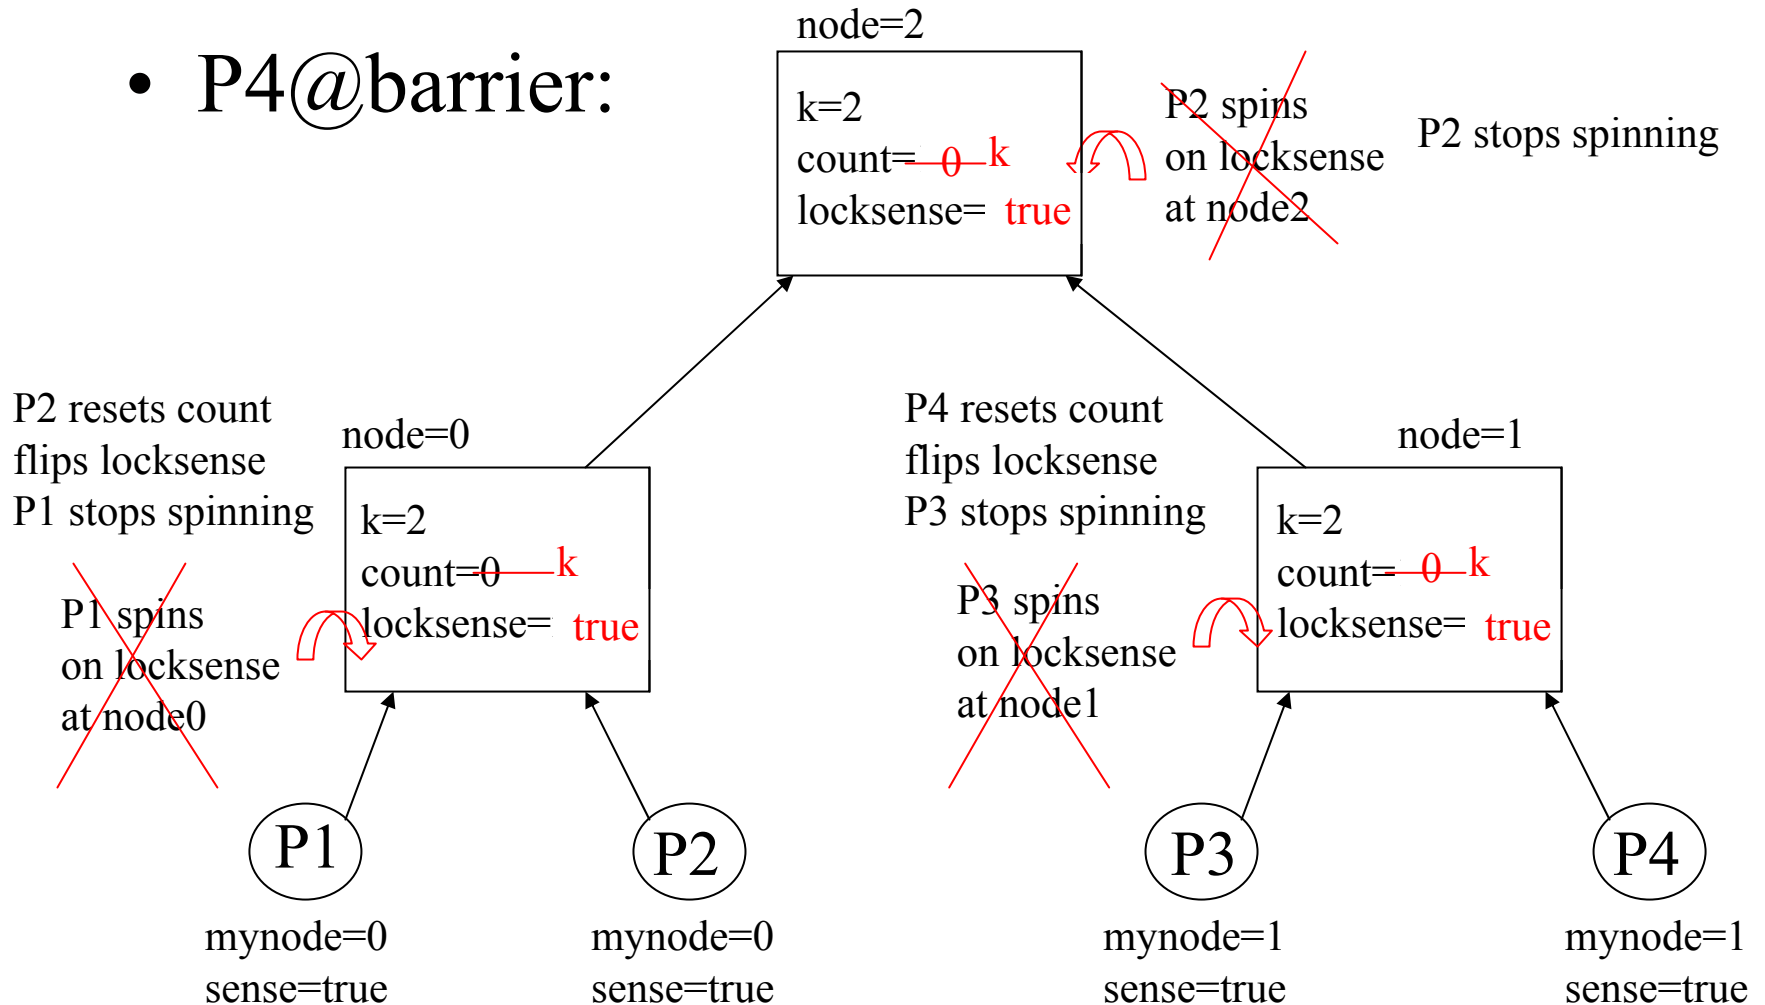
Software Combining Tree Barrier

- P3@barrier:



Software Combining Tree Barrier

- P4@barrier:



SCT barrier

- Good for other collectives too
- On CC arch., may spin locally, otherwise on remote locations
- Spin location not statically known – depends on arrival of processes – ultimately unbound # network transactions on nonCCwBroadcast
- Length of critical path
- Space

Other barrier versions in paper

- Dissemination:
 - Each proc notifies $\log P$ procs it reached a barrier, and also waits on $\log P$ procs
 - Complex/sizeable data structure $O(P \log P)$, simple signaling operations but $O(P \log P)$
 - Local spinning
- Tournament:
 - Like combining tree, but predetermined which processor goes up to the next level – statically known spinning locations
 - $O(P \log P)$ space, $O(P)$ network transactions
 - May spin on global flag for CCwBroadcast architectures
- Proposed barrier:
 - Again tree based, smaller data structure $O(P)$, and minimizes number of network transactions – still $O(P)$;
 - Spin optimization as above
- Each processor updates a dedicated (set of) locations – no need for atomic ops!




Figure 10: Performance of barriers on the Symmetry.




Figure 9: Performance of selected barriers on the Butterfly.